

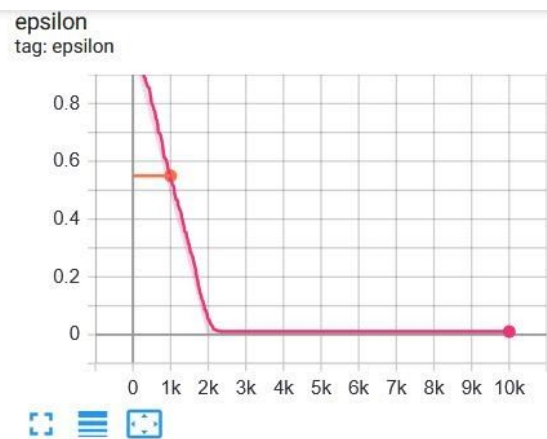
MT19213
Ghanendra Singh
DQN Simulation results.

1. Case. When numSteps = 10,000 and decaystop = 2000, ep = 0.01

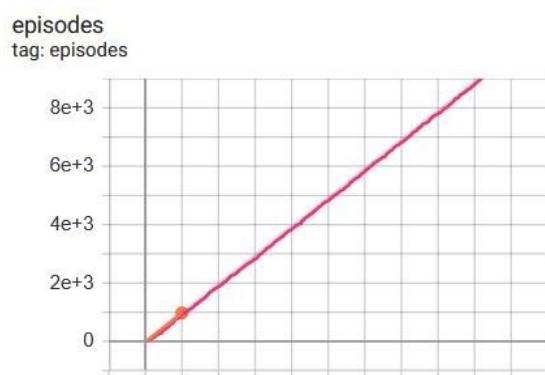
```
▶ trainer.train(numSteps=10000, decayStop=2000, endEpsilon=0.01)
```

Updated target network updates

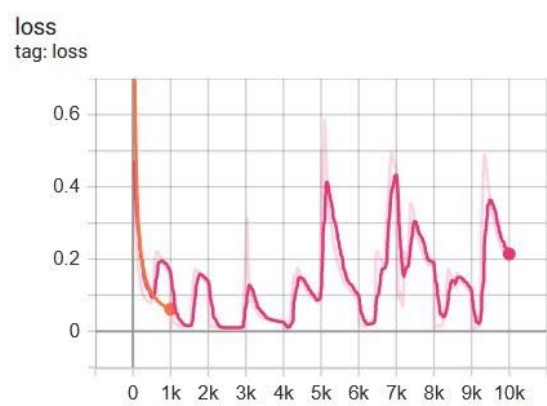
Tensorboard smoothing 0.9



episodes



loss



2. Case NumSteps = 10^5 , decayStop = 20K



```
trainer.train(numSteps=100000, decayStop=20000, endEpsilon=0.01)
```


```
Updated target network updates  
Updated target network updates  
Updated target network updates  
Updated target network updates  
Updated target network updates  
Updated target network updates  
Updated target network updates  
Updated target network updates  
Updated target network updates  
Updated target network updates
```

Tensorboard smoothing - 0.98

cum_reward
tag: cum_reward




epsilon
tag: epsilon



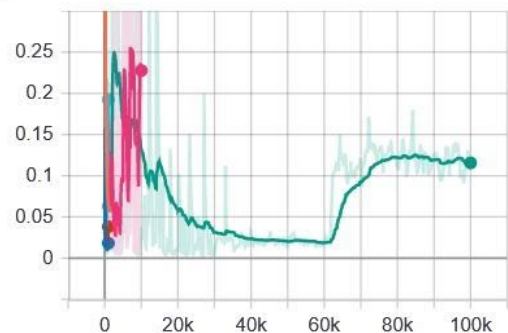
episodes
tag: episodes

episodes
tag: episodes



loss

loss
tag: loss

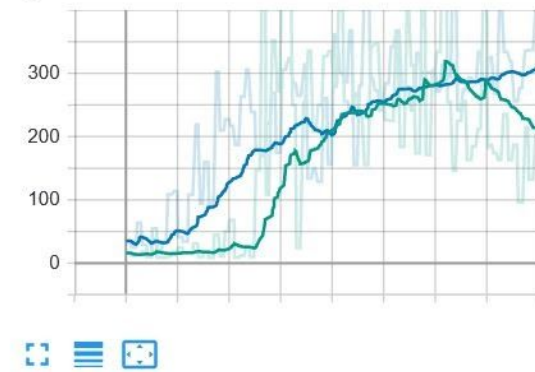


Q3. DDQN Simulation Results comparison with DQN.


When numSteps=100000, decayStop=20000

Tensorboard smoothing value 0.99 for removing the variations.

cum_reward
tag: cum_reward




epsilon
tag: epsilon




episodes

episodes
tag: episodes

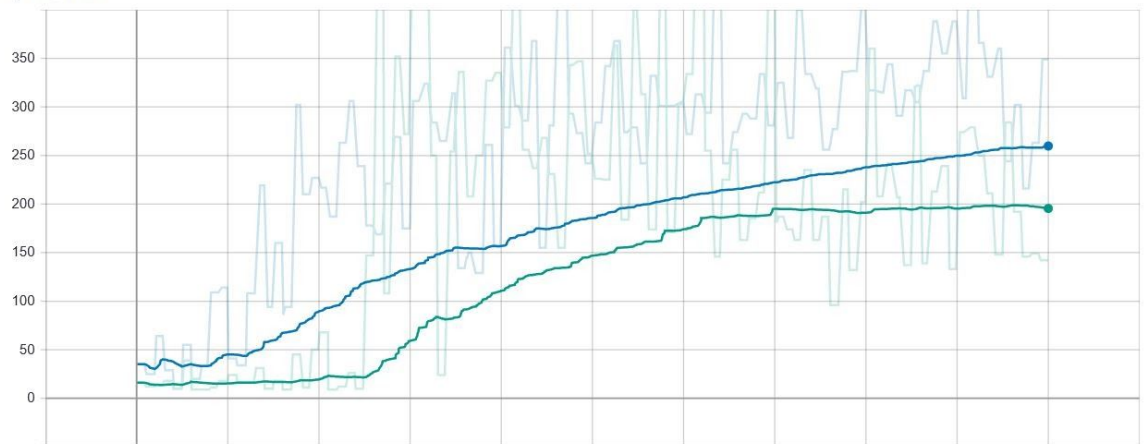


loss

loss
tag: loss



cum_reward
tag: cum_reward



Comparison.

1. Overestimates observed in the DQN case are reduced in DDQN case.
2. Here as the no. of steps are only 10^5 , less difference is visible in the estimate values.
3. Action selection and evaluation is done based on the same values in DQN where as in DDQN two value functions are learned in which instead of Q_a and Q_b , the rules were modified.

OLD Implementation

Algorithm 1 Double Q-learning

```
1: Initialize  $Q^A, Q^B, s$ 
2: repeat
3:   Choose  $a$ , based on  $Q^A(s, \cdot)$  and  $Q^B(s, \cdot)$ , observe  $r, s'$ 
4:   Choose (e.g. random) either UPDATE(A) or UPDATE(B)
5:   if UPDATE(A) then
6:     Define  $a^* = \arg \max_a Q^A(s', a)$ 
7:      $Q^A(s, a) \leftarrow Q^A(s, a) + \alpha(s, a) (r + \gamma Q^B(s', a^*) - Q^A(s, a))$ 
8:   else if UPDATE(B) then
9:     Define  $b^* = \arg \max_a Q^B(s', a)$ 
10:     $Q^B(s, a) \leftarrow Q^B(s, a) + \alpha(s, a) (r + \gamma Q^A(s', b^*) - Q^B(s, a))$ 
11:   end if
12:    $s \leftarrow s'$ 
13: until end
```

Pseudo-code Source: "Double Q-learning" (Hasselt, 2010)

NEW Implementation

Algorithm 1 : Double Q-learning (Hasselt et al., 2015)

```
Initialize primary network  $Q_\theta$ , target network  $Q_{\theta'}$ , replay buffer  $\mathcal{D}$ ,  $\tau \ll 1$ 
for each iteration do
  for each environment step do
    Observe state  $s_t$  and select  $a_t \sim \pi(a_t, s_t)$ 
    Execute  $a_t$  and observe next state  $s_{t+1}$  and reward  $r_t = R(s_t, a_t)$ 
    Store  $(s_t, a_t, r_t, s_{t+1})$  in replay buffer  $\mathcal{D}$ 
  for each update step do
    sample  $e_t = (s_t, a_t, r_t, s_{t+1}) \sim \mathcal{D}$ 
    Compute target Q value:
       $Q^*(s_t, a_t) \approx r_t + \gamma Q_{\theta'}(s_{t+1}, \arg \max_{a'} Q_{\theta'}(s_{t+1}, a'))$ 
    Perform gradient descent step on  $(Q^*(s_t, a_t) - Q_\theta(s_t, a_t))^2$ 
    Update target network parameters:
       $\theta' \leftarrow \tau * \theta + (1 - \tau) * \theta'$ 
```
